

# Data Cleaning Workflow Template

This notebook is a reusable template for cleaning raw data:

1. Load data
2. Inspect structure
3. Handle missing values
4. Remove duplicates
5. Fix data types
6. Clean text/categories
7. Detect & cap outliers (IQR)
8. Validate & export

```
In [11]: # === 1. IMPORTS & SETTINGS ===
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

pd.set_option('display.max_columns', None)
pd.set_option('display.float_format', lambda x: f'{x:.2f}')

print("Libraries loaded.")
```

Libraries loaded.

```
In [6]: # === 2. LOAD RAW DATA ===
```

```
# TODO: Change the file name/path for each new project.
```

```
DATA_PATH = "/home/jkatz015/repos/personal/portfolio/projects/project-02-ap-automation/data/invoices_raw/ap_transactions.csv"
df = pd.read_csv(DATA_PATH)

print("Rows:", len(df))
df.head()
```

Rows: 50150

Out[6]:

|   | INVOICE_ID | VENDOR_ID | VENDOR_NAME          | VENDOR_NUMBER | VENDOR_COUNTRY    | VENDOR_REGION | INVOICE_DATE | DU_DATE    |
|---|------------|-----------|----------------------|---------------|-------------------|---------------|--------------|------------|
| 0 | 873641333  | 846210656 | Metro Solutions      | 1230020547    | Brunei Darussalam | North-Region  | 2022-04-11   | 2022-04-15 |
| 1 | 865186916  | 892667166 | Metro Corp           | 1230022657    | Spain             | North-Region  | 2022-03-09   | 2022-03-13 |
| 2 | 914629138  | 921212744 | Superior Enterprises | 1230002613    | Kenya             | South-Region  | 2022-02-23   | 2022-02-28 |
| 3 | 861070279  | 803387626 | Global Manufacturing | 1230014342    | Turkmenistan      | North-Region  | 2022-02-01   | 2022-02-05 |
| 4 | 784182066  | 942636818 | Quality Partners     | 1230014342    | Guam              | North-Region  | 2022-02-14   | 2022-02-18 |



In [12]:

# === 3. BASIC STRUCTURE CHECK ===

```
df.info()  
df.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50150 entries, 0 to 50149
Data columns (total 20 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   INVOICE_ID       50150 non-null   int64  
 1   VENDOR_ID        50150 non-null   int64  
 2   VENDOR_NAME      50150 non-null   object  
 3   VENDOR_NUMBER    50150 non-null   int64  
 4   VENDOR_COUNTRY   50150 non-null   object  
 5   VENDOR_REGION    50150 non-null   object  
 6   INVOICE_DATE     50150 non-null   object  
 7   DUE_DATE         50150 non-null   object  
 8   INVOICE_AMOUNT   50150 non-null   float64 
 9   CURRENCY         50150 non-null   object  
 10  GL_CODE          50150 non-null   object  
 11  PO_TYPE          50150 non-null   object  
 12  COMPANY_CODE     50150 non-null   int64  
 13  BUSINESS_UNIT    50150 non-null   int64  
 14  COST_CENTER      50150 non-null   object  
 15  PAYMENT_STATUS   50150 non-null   object  
 16  APPROVAL_STATUS  50150 non-null   float64 
 17  CREDIT_LIMIT     50150 non-null   float64 
 18  INVOICE_TIME     50150 non-null   int64  
 19  _ANOMALY_TYPE    50150 non-null   object  
dtypes: float64(3), int64(6), object(11)
memory usage: 7.7+ MB
```

Out[12]:

|              | INVOICE_ID    | VENDOR_ID    | VENDOR_NUMBER | INVOICE_AMOUNT | COMPANY_CODE | BUSINESS_UNIT | APPROVAL_STATI |
|--------------|---------------|--------------|---------------|----------------|--------------|---------------|----------------|
| <b>count</b> | 50150.00      | 50150.00     | 50150.00      | 50150.00       | 50150.00     | 50150.00      | 50150.         |
| <b>mean</b>  | 879640407.81  | 876513993.64 | 1111686783.56 | 6317.53        | 3415.36      | 3304.16       | 66.            |
| <b>std</b>   | 86465713.92   | 70885396.17  | 357815061.23  | 216296.24      | 594.99       | 696.53        | 8.             |
| <b>min</b>   | 753460934.00  | 753458315.00 | 12100061.00   | 0.01           | 59.00        | 2100.00       | 52.            |
| <b>25%</b>   | 815645178.75  | 815235731.75 | 1230011793.00 | 189.04         | 3260.00      | 2702.00       | 64.            |
| <b>50%</b>   | 877306681.50  | 876551085.50 | 1230014314.00 | 457.02         | 3660.00      | 3300.00       | 64.            |
| <b>75%</b>   | 938984852.25  | 938060359.25 | 1230014342.00 | 1057.20        | 3660.00      | 3908.00       | 64.            |
| <b>max</b>   | 1897434115.00 | 999007452.00 | 1230025103.00 | 44854887.00    | 4260.00      | 4510.00       | 93.            |



In [8]: # === 4. MISSING VALUES OVERVIEW ===

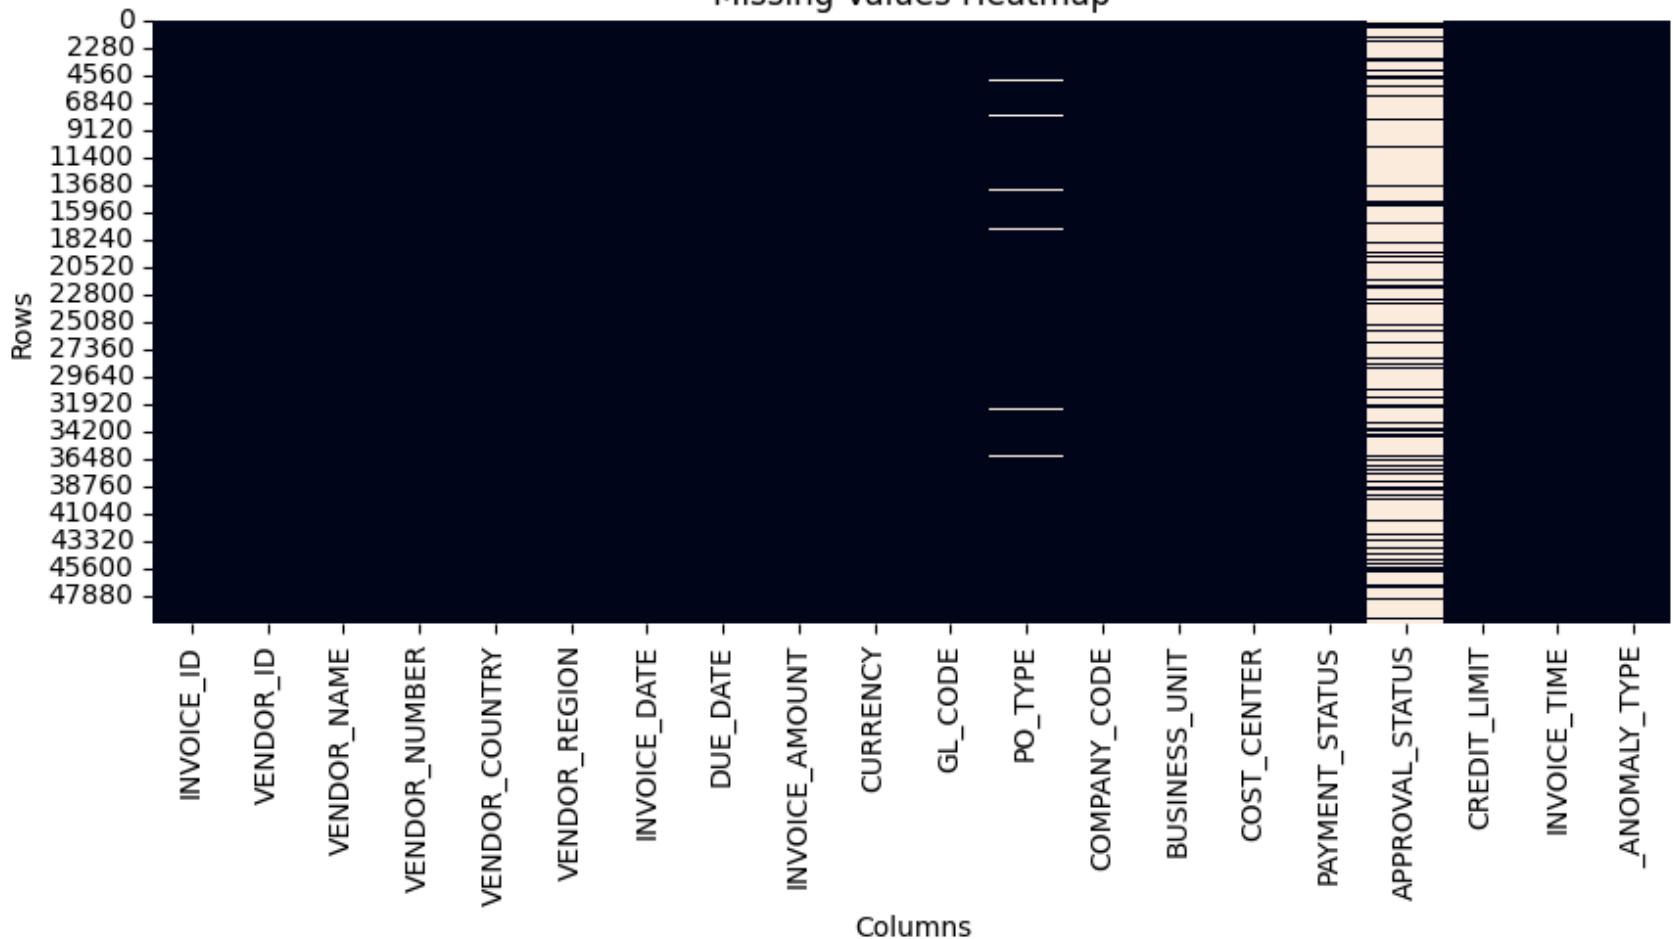
```
print("Missing values per column:")
print(df.isnull().sum())

plt.figure(figsize=(10, 4))
sns.heatmap(df.isnull(), cbar=False)
plt.title("Missing Values Heatmap")
plt.xlabel("Columns")
plt.ylabel("Rows")
plt.show()
```

Missing values per column:

```
INVOICE_ID          0
VENDOR_ID          0
VENDOR_NAME        0
VENDOR_NUMBER      0
VENDOR_COUNTRY     0
VENDOR_REGION      0
INVOICE_DATE        0
DUE_DATE           0
INVOICE_AMOUNT      0
CURRENCY           0
GL_CODE            0
PO_TYPE            1063
COMPANY_CODE       0
BUSINESS_UNIT      0
COST_CENTER        0
PAYMENT_STATUS     0
APPROVAL_STATUS    38811
CREDIT_LIMIT       0
INVOICE_TIME        0
_ANOMALY_TYPE      0
dtype: int64
```

## Missing Values Heatmap



```
In [13]: # === 5. IDENTIFY COLUMN TYPES ===

numeric_cols = df.select_dtypes(include=['float64', 'int64', 'Int64']).columns.tolist()
text_cols    = df.select_dtypes(include=['object']).columns.tolist()

print("Numeric columns:", numeric_cols)
print("Text columns:", text_cols)
```

```
Numeric columns: ['INVOICE_ID', 'VENDOR_ID', 'VENDOR_NUMBER', 'INVOICE_AMOUNT', 'COMPANY_CODE', 'BUSINESS_UNIT', 'APP  
ROVAL_STATUS', 'CREDIT_LIMIT', 'INVOICE_TIME']  
Text columns: ['VENDOR_NAME', 'VENDOR_COUNTRY', 'VENDOR_REGION', 'INVOICE_DATE', 'DUE_DATE', 'CURRENCY', 'GL_CODE',  
'PO_TYPE', 'COST_CENTER', 'PAYMENT_STATUS', '_ANOMALY_TYPE']
```

```
In [14]: # === 6. HANDLE MISSING VALUES ===  
# Adjust per project if needed.  
  
# 6a. Numeric → fill with median  
for col in numeric_cols:  
    if df[col].isnull().sum() > 0:  
        median_value = df[col].median()  
        df.loc[:, col] = df[col].fillna(median_value)  
  
# 6b. Text → fill with "Unknown"  
for col in text_cols:  
    if df[col].isnull().sum() > 0:  
        df.loc[:, col] = df[col].fillna("Unknown")  
  
print("Missing values after fill:")  
print(df.isnull().sum())
```

```
Missing values after fill:  
INVOICE_ID      0  
VENDOR_ID       0  
VENDOR_NAME     0  
VENDOR_NUMBER   0  
VENDOR_COUNTRY  0  
VENDOR_REGION   0  
INVOICE_DATE    0  
DUE_DATE        0  
INVOICE_AMOUNT   0  
CURRENCY         0  
GL_CODE          0  
PO_TYPE          0  
COMPANY_CODE    0  
BUSINESS_UNIT    0  
COST_CENTER     0  
PAYMENT_STATUS   0  
APPROVAL_STATUS  0  
CREDIT_LIMIT    0  
INVOICE_TIME     0  
_ANOMALY_TYPE    0  
dtype: int64
```

```
In [15]: # === 7. REMOVE DUPLICATES ===  
  
dups_before = df.duplicated().sum()  
print("Duplicates before:", dups_before)  
  
df = df.drop_duplicates()  
  
dups_after = df.duplicated().sum()  
print("Duplicates after:", dups_after)
```

```
Duplicates before: 0  
Duplicates after: 0
```

```
In [16]: # === 8. FIX DATA TYPES ===  
  
# Example: convert columns whose name suggests date/time  
date_like_cols = [c for c in df.columns if 'date' in c.lower()]  
for col in date_like_cols:  
    df.loc[:, col] = pd.to_datetime(df[col], errors='coerce')
```

```
# Example: convert typical integer-like columns if present
int_candidate_cols = ['age', 'quantity', 'count'] # edit per dataset
for col in int_candidate_cols:
    if col in df.columns:
        df.loc[:, col] = df[col].astype('Int64')

print(df.dtypes)
```

```
INVOICE_ID      int64
VENDOR_ID       int64
VENDOR_NAME     object
VENDOR_NUMBER   int64
VENDOR_COUNTRY  object
VENDOR_REGION   object
INVOICE_DATE    object
DUE_DATE        object
INVOICE_AMOUNT  float64
CURRENCY        object
GL_CODE          object
PO_TYPE          object
COMPANY_CODE    int64
BUSINESS_UNIT   int64
COST_CENTER     object
PAYMENT_STATUS  object
APPROVAL_STATUS float64
CREDIT_LIMIT    float64
INVOICE_TIME    int64
_ANOMALY_TYPE   object
dtype: object
```

```
In [19]: # === 9. CLEAN TEXT / CATEGORICAL COLUMNS ===
```

```
for col in text_cols:
    df.loc[:, col] = df[col].astype(str)
    df.loc[:, col] = df[col].str.strip()
    df.loc[:, col] = df[col].str.upper()

# Example: inspect a key category, if it exists
if 'CITY' in df.columns:
    print(df['CITY'].value_counts().head(20))
```

```
In [20]: # === 10. OUTLIER DETECTION & CAPPING (IQR) ===
```

```
def iqr_bounds(series: pd.Series):
    Q1 = series.quantile(0.25)
    Q3 = series.quantile(0.75)
    IQR = Q3 - Q1
    lower = Q1 - 1.5 * IQR
    upper = Q3 + 1.5 * IQR
    return lower, upper

for col in numeric_cols:
    if df[col].nunique() <= 1:
        continue # skip constant columns

    lower, upper = iqr_bounds(df[col])
    print(f"{col}: lower={lower:.2f}, upper={upper:.2f}")

    df.loc[df[col] < lower, col] = lower
    df.loc[df[col] > upper, col] = upper
```

```
INVOICE_ID: lower=630635668.50, upper=1123994362.50
VENDOR_ID: lower=630998790.50, upper=1122297300.50
VENDOR_NUMBER: lower=1230007969.50, upper=1230018165.50
INVOICE_AMOUNT: lower=-1113.21, upper=2359.44
COMPANY_CODE: lower=2660.00, upper=4260.00
BUSINESS_UNIT: lower=893.00, upper=5717.00
INVOICE_TIME: lower=-67351.12, upper=337595.88
```

```
In [21]: # === 11. FINAL VALIDATION ===
```

```
print("Missing values per column:")
print(df.isnull().sum())

print("\nDuplicates:", df.duplicated().sum())

print("\nData types:")
print(df.dtypes)

print("\nPreview:")
df.head()
```

Missing values per column:

```
INVOICE_ID      0
VENDOR_ID       0
VENDOR_NAME     0
VENDOR_NUMBER   0
VENDOR_COUNTRY  0
VENDOR_REGION   0
INVOICE_DATE    0
DUE_DATE        0
INVOICE_AMOUNT   0
CURRENCY         0
GL_CODE          0
PO_TYPE          0
COMPANY_CODE    0
BUSINESS_UNIT   0
COST_CENTER     0
PAYMENT_STATUS   0
APPROVAL_STATUS  0
CREDIT_LIMIT    0
INVOICE_TIME    0
_ANOMALY_TYPE   0
dtype: int64
```

Duplicates: 0

Data types:

```
INVOICE_ID      float64
VENDOR_ID       float64
VENDOR_NAME     object
VENDOR_NUMBER   float64
VENDOR_COUNTRY  object
VENDOR_REGION   object
INVOICE_DATE    object
DUE_DATE        object
INVOICE_AMOUNT   float64
CURRENCY        object
GL_CODE          object
PO_TYPE          object
COMPANY_CODE    int64
BUSINESS_UNIT   int64
COST_CENTER     object
PAYMENT_STATUS   object
```

```
APPROVAL_STATUS      float64  
CREDIT_LIMIT         float64  
INVOICE_TIME          float64  
_ANOMALY_TYPE        object  
dtype: object
```

Preview:

Out[21]:

|   | INVOICE_ID   | VENDOR_ID    | VENDOR_NAME          | VENDOR_NUMBER | VENDOR_COUNTRY    | VENDOR_REGION | INVOICE_DATE        | DUE_DATE            |
|---|--------------|--------------|----------------------|---------------|-------------------|---------------|---------------------|---------------------|
| 0 | 873641333.00 | 846210656.00 | METRO SOLUTIONS      | 1230018165.50 | BRUNEI DARUSSALAM | NORTH-REGION  | 2022-04-11 00:00:00 | 2022-04-15 00:00:00 |
| 1 | 865186916.00 | 892667166.00 | METRO CORP           | 1230018165.50 | SPAIN             | NORTH-REGION  | 2022-03-09 00:00:00 | 2022-03-13 00:00:00 |
| 2 | 914629138.00 | 921212744.00 | SUPERIOR ENTERPRISES | 1230007969.50 | KENYA             | SOUTH-REGION  | 2022-02-23 00:00:00 | 2022-02-27 00:00:00 |
| 3 | 861070279.00 | 803387626.00 | GLOBAL MANUFACTURING | 1230014342.00 | TURKMENISTAN      | NORTH-REGION  | 2022-02-01 00:00:00 | 2022-02-05 00:00:00 |
| 4 | 784182066.00 | 942636818.00 | QUALITY PARTNERS     | 1230014342.00 | GUAM              | NORTH-REGION  | 2022-02-14 00:00:00 | 2022-02-18 00:00:00 |



In [ ]: